

# Introduction to Machine Learning

## Chap II: Regression

Julien Donini

LPC/Université Clermont Auvergne

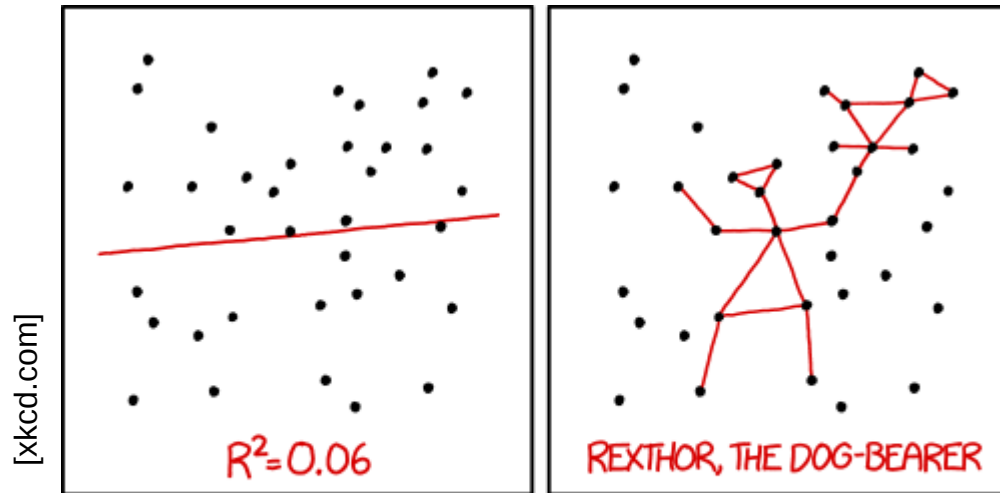
## Linear regression

- 1D and multidimensional data
- Basis functions
- Example: polynomial curve fitting
- Regularization (\*)

## Likelihood and regression (\*)

- Variable/feature:  $x$ , weight:  $w$
- Vector: variables  $\mathbf{x} = \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$ ; weights  $\mathbf{w} = \begin{pmatrix} w_1 \\ \vdots \\ w_n \end{pmatrix}$
- Vector (transpose):  $\mathbf{x}^T = (x_1, \dots, x_n)$
- Dot product :  $\mathbf{w} \cdot \mathbf{x} = \mathbf{w}^T \mathbf{x} = \sum_{i=1}^n w_i x_i = w_1 x_1 + w_2 x_2 + \dots w_n x_n.$
- Sequence of  $p$  vectors:  $\{\mathbf{x}_j\}_{(j=1..p)}$
- Matrix (size  $n \times m$ ):  $\mathbf{M} = \begin{pmatrix} a_{11} & \cdots & a_{1m} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nm} \end{pmatrix}$

# Linear regression



I DON'T TRUST LINEAR REGRESSIONS WHEN IT'S HARDER  
TO GUESS THE DIRECTION OF THE CORRELATION FROM THE  
SCATTER PLOT THAN TO FIND NEW CONSTELLATIONS ON IT.

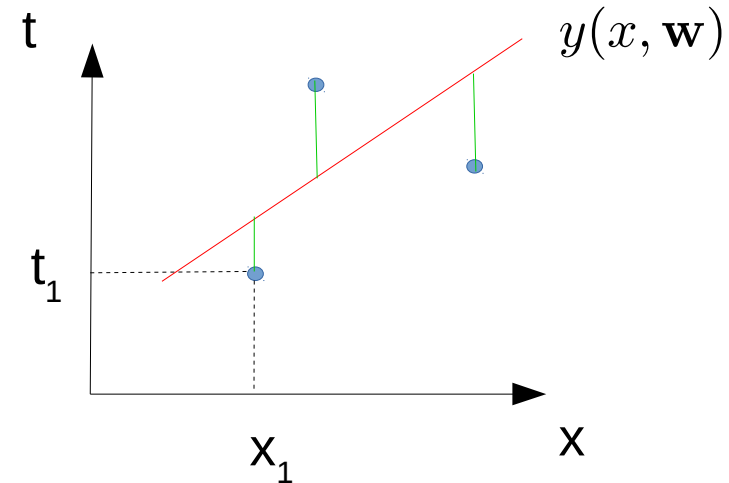
# Simple case: 1 dimensional data

## Training dataset

- N observations of **feature**  $x = \{x_1, \dots, x_N\}$
- N **Target** values  $t = \{t_1, \dots, t_N\}$

**Prediction model:** straight line

$$y(x, \mathbf{w}) = y(x; w_0, w_1) = w_0 + w_1 x$$



**Weights** determined by minimizing an **Error function E**

- also called **Cost function** or **Loss function** (i.e what is the consequence of your error !)

Common choice: sum of **square distance** between function and target:

$$E(w_0, w_1) = \sum_{i=1}^N \{y(x_i; w_0, w_1) - t_i\}^2$$

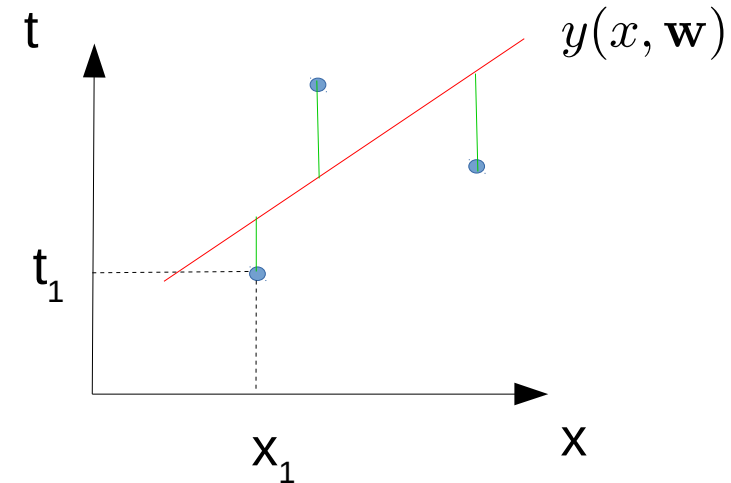
# Simple case: 1 dimensional data

## Training dataset

- N observations of **feature**  $x = \{x_1, \dots, x_N\}$
- N **Target** values  $t = \{t_1, \dots, t_N\}$

## Prediction model: straight line

$$y(x, \mathbf{w}) = y(x; w_0, w_1) = w_0 + w_1 x$$



Here **optimal weights** can be calculated analytically (not always possible !)

$$E(w_0, w_1) = \sum_{i=1}^N \{y(x_i; w_0, w_1) - t_i\}^2$$
$$\begin{cases} \frac{\partial E(w_0, w_1)}{\partial w_0} = 0 \\ \frac{\partial E(w_0, w_1)}{\partial w_1} = 0 \end{cases} \Leftrightarrow \begin{cases} w_1 = \frac{\text{cov}(x, t)}{\text{var}(x)} = r \frac{\sigma(t)}{\sigma(x)} \\ w_0 = \bar{t} - r \frac{\sigma(t)}{\sigma(x)} \bar{x} \end{cases}$$

(r: correlation factor between x and t)

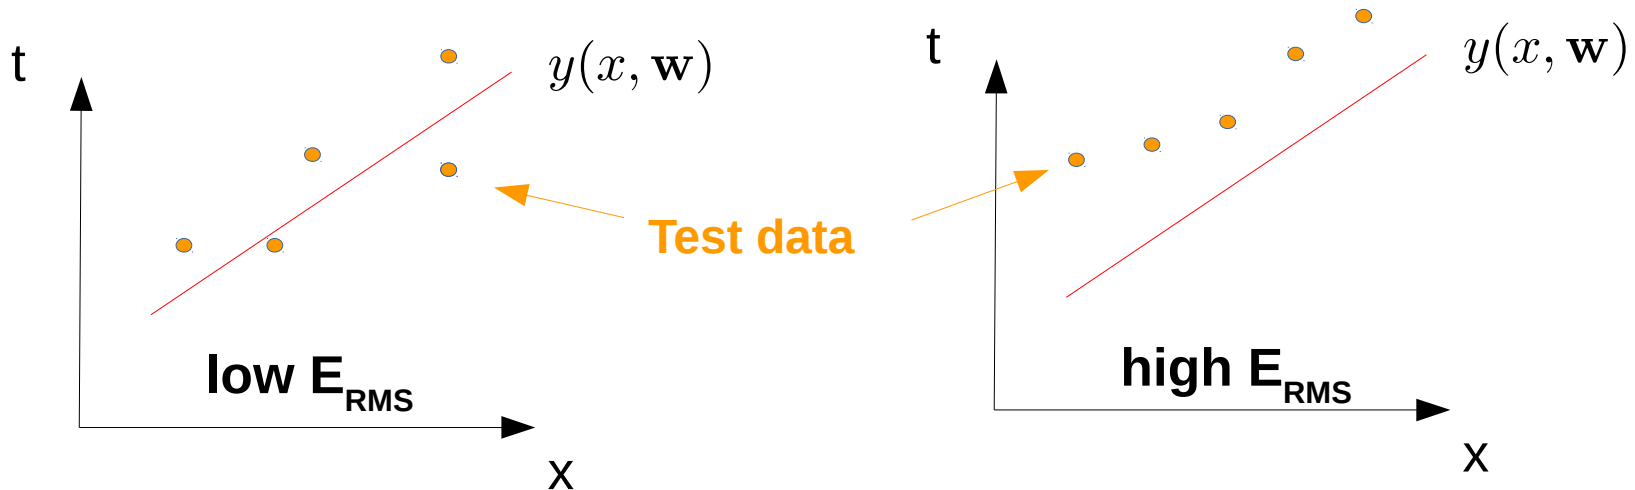
# Simple case: 1 dimensional data

## Training and testing

- **Training**: use dataset to determine **weights**  $w_0$  and  $w_1$
- **Testing**: check compatibility of  $y(x, \mathbf{w})$  on a new dataset

Measure of **compatibility**: root mean squared error (RMS)

$$E_{RMS} = \sqrt{\frac{1}{N} \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2} = \sqrt{\frac{E(\mathbf{w})}{N}}$$



# Generalization: multidimensional data

## Dataset ( $p \times 1$ data)

- $N$  observations of  $p$ -dimensions features

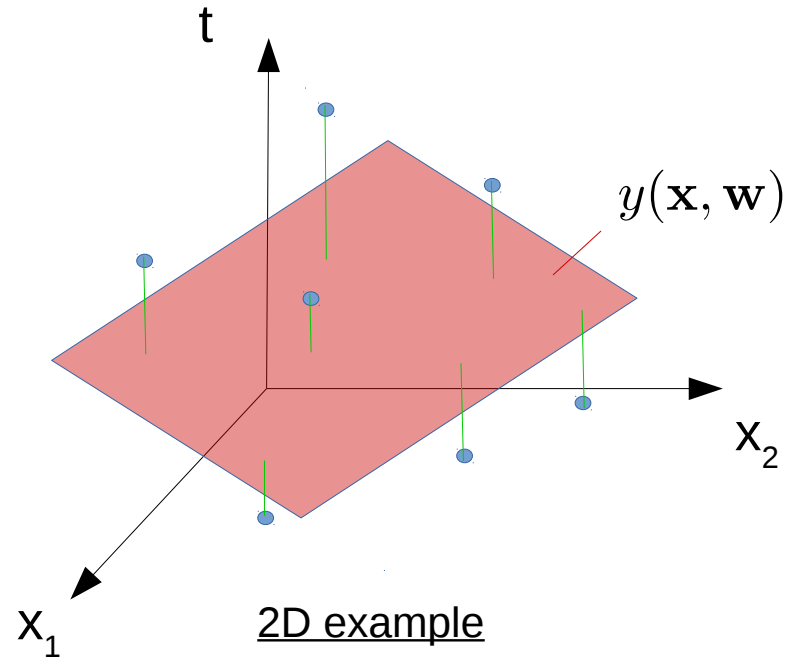
$$\{\mathbf{x}_i\}_{i=1..N} = \{\mathbb{R}^p\} = \left\{ \begin{pmatrix} x_1 \\ \vdots \\ x_p \end{pmatrix} \right\}$$

- $N$  target values  $\mathbf{t} = \{t_1, \dots, t_N\}$

## Fit function: multidimensional plane

- Linear function with  $p+1$  weights:  $\mathbf{w}$

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \mathbf{w}^T \mathbf{x} = \underbrace{w_0}_{\text{bias term}} + w_1 x_1 + w_2 x_2 + \dots w_p x_p.$$





# Generalization: multidimensional data

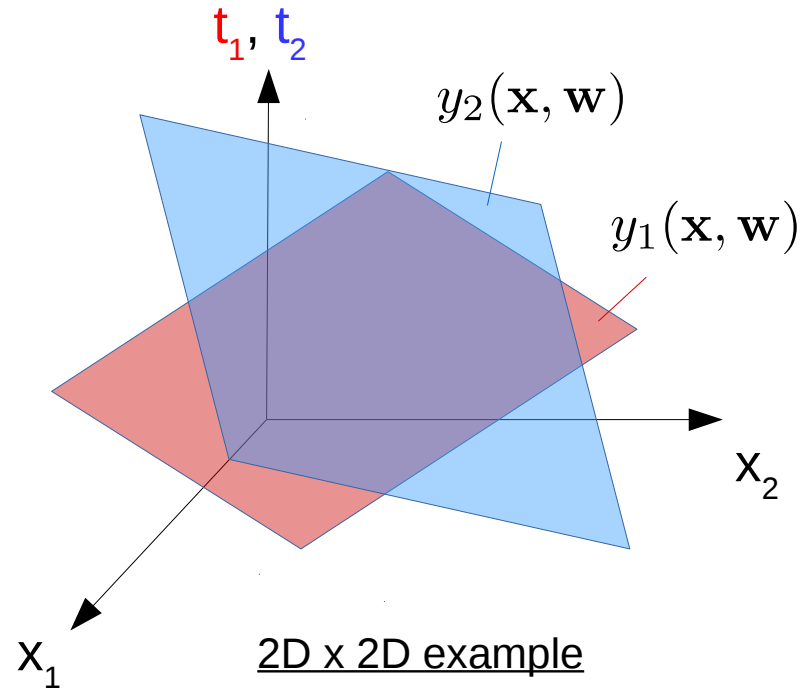
## Dataset (p x q data)

- **N** observations of **p**-dimensions **features**

$$\{\mathbf{x}_i\}_{i=1..N} = \{\mathbb{R}^p\} = \left\{ \begin{pmatrix} x_1 \\ \vdots \\ x_p \end{pmatrix} \right\}$$

- **N target** values of **q**-dimensions

$$\{\mathbf{t}_i\}_{i=1..N} = \{\mathbb{R}^q\} = \left\{ \begin{pmatrix} t_1 \\ \vdots \\ t_q \end{pmatrix} \right\}$$



## Fit functions:

$$\begin{pmatrix} y_1(\mathbf{x}, \mathbf{w}) \\ \vdots \\ y_q(\mathbf{x}, \mathbf{w}) \end{pmatrix} = \underbrace{\begin{pmatrix} w_{01} \\ \vdots \\ w_{0q} \end{pmatrix}}_{\text{bias terms}} + \begin{pmatrix} w_{11} & \cdots & w_{1p} \\ \vdots & \ddots & \vdots \\ w_{q1} & \cdots & w_{qp} \end{pmatrix} \begin{pmatrix} x_1 \\ \vdots \\ x_p \end{pmatrix}$$

# Linear basis function models

Apply **M non-linear basis functions**  $\phi$  to input feature  $\mathbf{x}$ :

$$\mathbf{x} \longrightarrow \begin{pmatrix} \phi_1(\mathbf{x}) \\ \vdots \\ \phi_M(\mathbf{x}) \end{pmatrix} \quad \phi_j(\mathbf{x}): \text{basis function}$$

The regression function  $y(\mathbf{x}, \mathbf{w})$  then become non-linear function of  $\mathbf{x}$ :

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{i=1}^M w_i \phi_i(\mathbf{x}) = w_0 + w_1 \phi_1(\mathbf{x}) + \cdots + w_M \phi_M(\mathbf{x})$$

These functions are called **linear models** because they are linear in  $\mathbf{w}$ .

For high number of dimensions linear models suffer from **limitations**, and other approaches (as Neural Networks) are more suited.

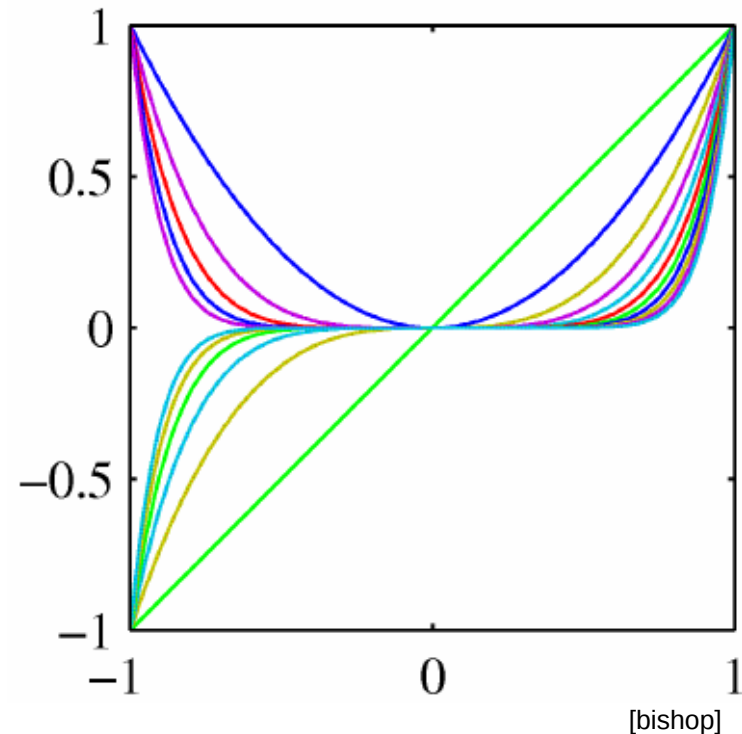
# Linear basis function models

## Polynomial basis functions (1D)

$$\phi_j(x) = x^j$$

$$y(x, \mathbf{w}) = \sum_{j=0}^{M-1} w_j x^j$$

**Global** functions of input variable  
→ a small change in  $x$  affects all  
basis functions



# Linear basis function models

## Gaussian basis functions (1D)

$$\phi_j(x) = e^{-\frac{(x-\mu_j)^2}{2\sigma^2}}$$

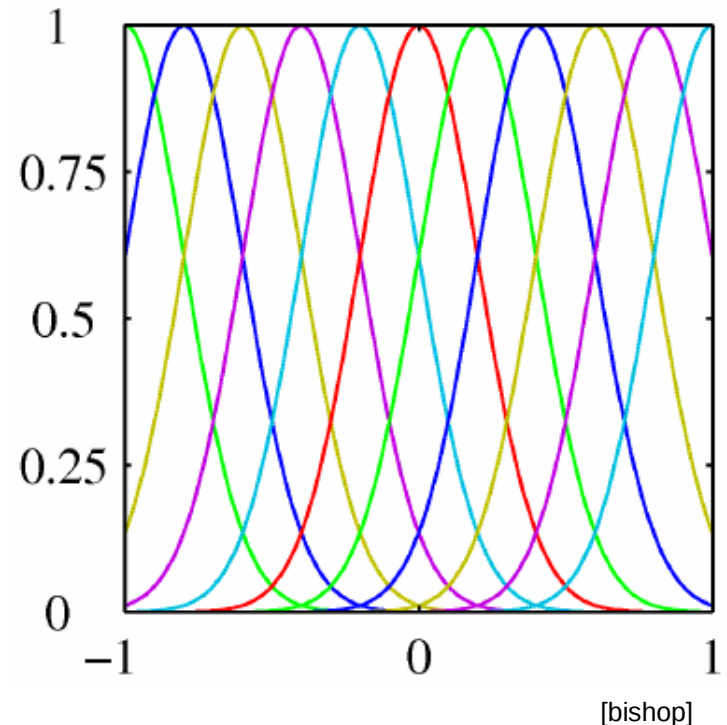
$$y(x, \mathbf{w}) = \sum_{j=0}^{M-1} w_j e^{-\frac{(x-\mu_j)^2}{2\sigma^2}}$$

### Parameters:

$\mu_j$  (location) and  $\sigma$  (width)

Normalization is not relevant.

**local** functions of input variable  
→ a small change in  $x$  mostly  
affects nearby basis functions



# Linear basis function models

## Sigmoidal basis functions (1D)

$$\phi_j(x) = \sigma \left( \frac{(x - \mu_j)}{s} \right)$$

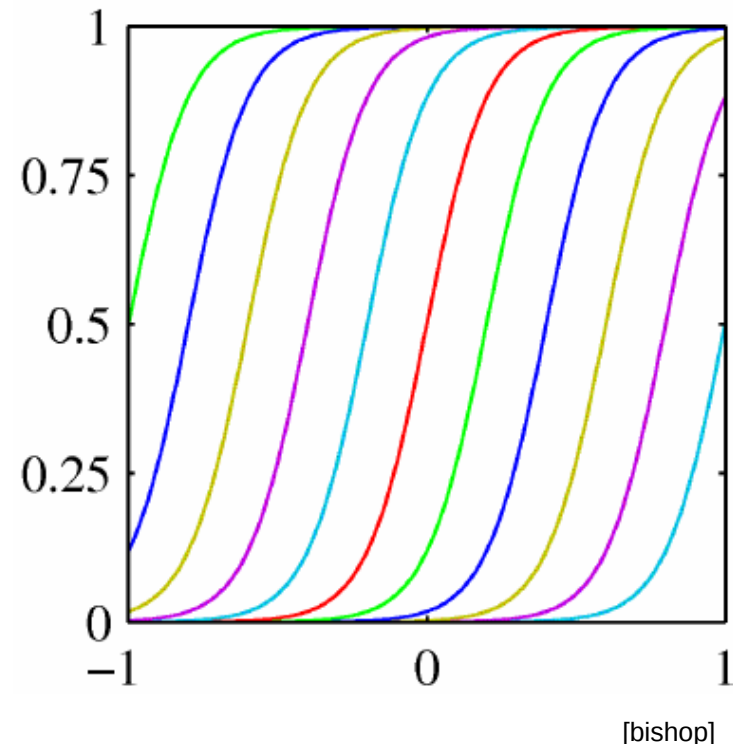
with

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

**Parameters:**

$\mu_j$  (location) and  $s$  (slope)

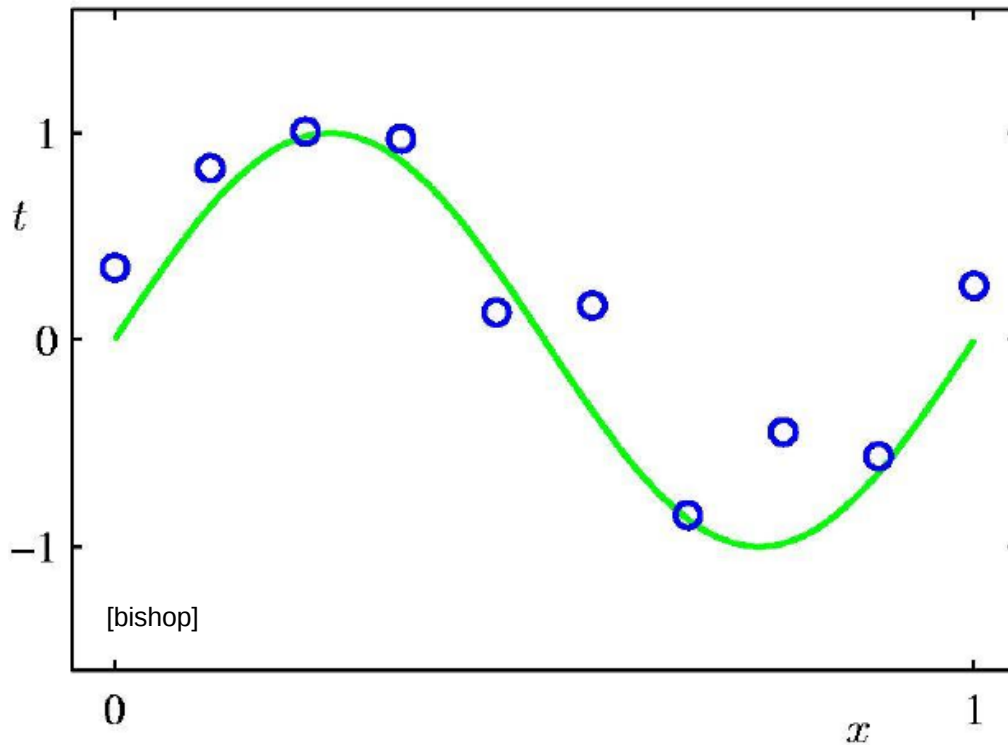
**local** functions of input variable  
→ a small change in  $x$  mostly  
affects nearby basis functions



# Example: polynomial curve fitting

## Training dataset

- $N$  observations of  $x = \{x_1, \dots, x_N\}$ : uniformly spaced in  $[0,1]$
- Target values  $t = \{t_1, \dots, t_N\}$ :  $\sin(2\pi x) + \text{Gaussian noise}$



Dummy example but could be e.g. temperature ( $t$ ) evolution over 1 day ( $x$ )

# Polynomial curve fitting

## Fit function

- Polynomial function of degree **M**, with coefficients  $\mathbf{w} = (w_1, \dots, w_M)^\top$

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

- Non-linear function of  $x$ , but linear function of  $\mathbf{w}$  → **linear model**
- Values of coefficient obtained by **minimizing** an **error function**
- Sum of the square of the errors**  $E(\mathbf{w})$

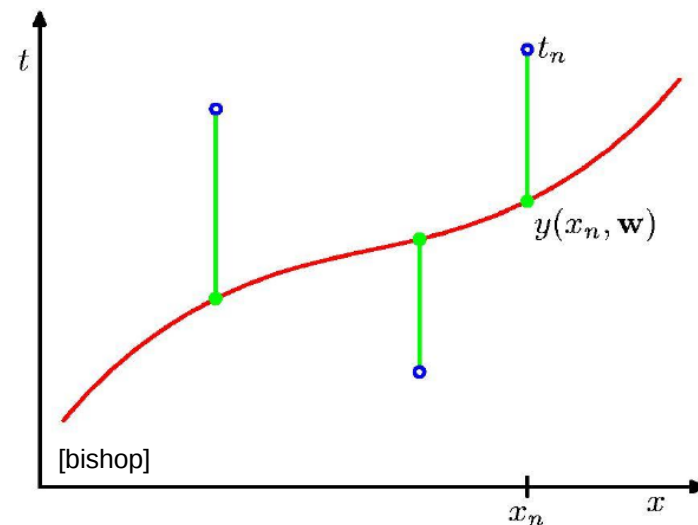
$$E(\mathbf{w}) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2$$

Minimization

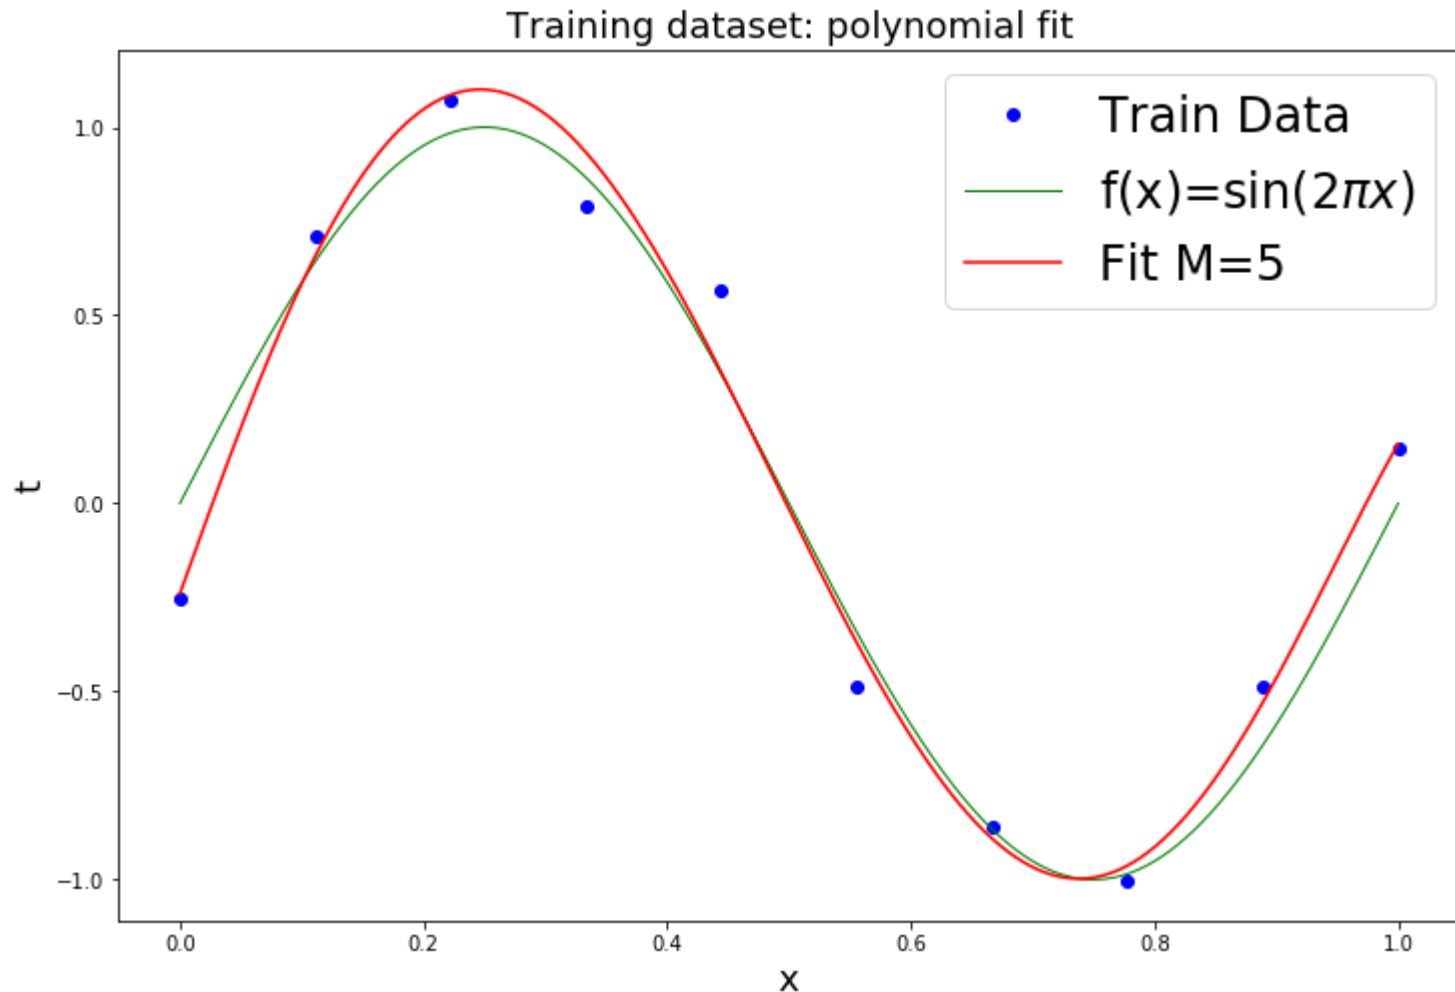


Fitted weights  $\mathbf{w}^*$

$E(\mathbf{w}^*)$



# Tutorial: python code



<https://mybinder.org/v2/gh/judonini/MLcourses/master>

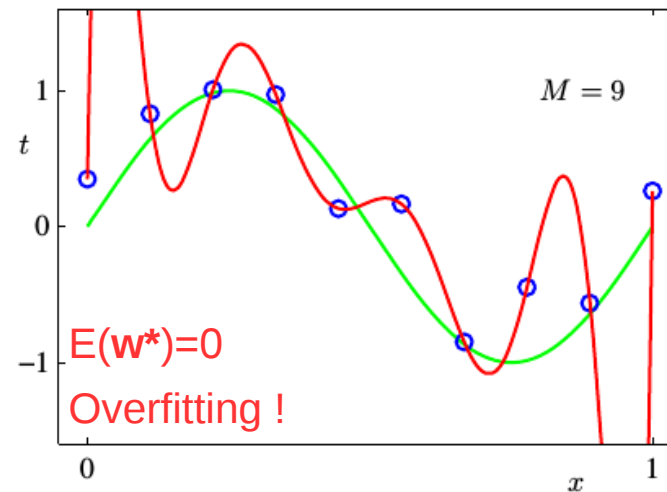
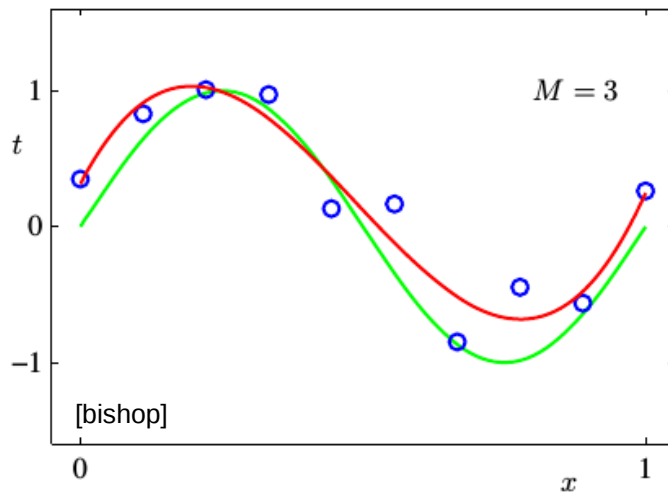
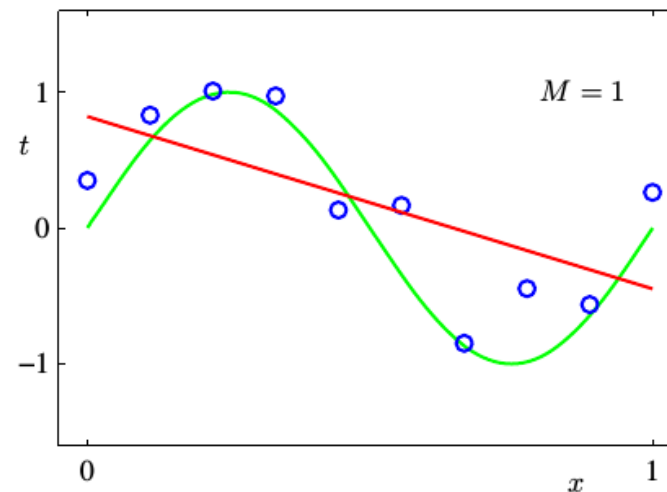
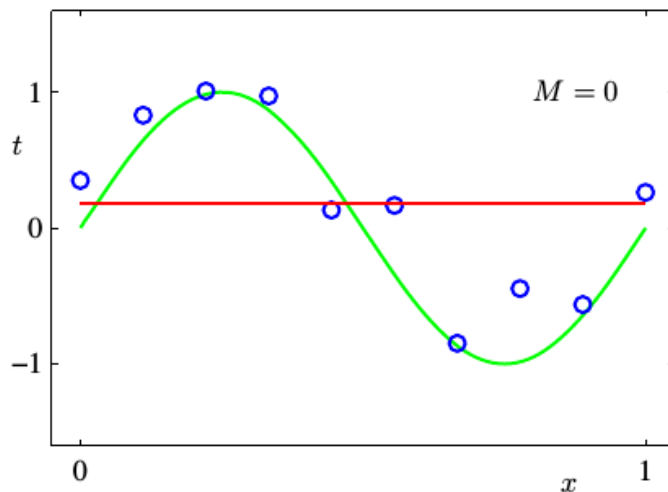


→ Polynomial-regression.ipynb



# Overfitting

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

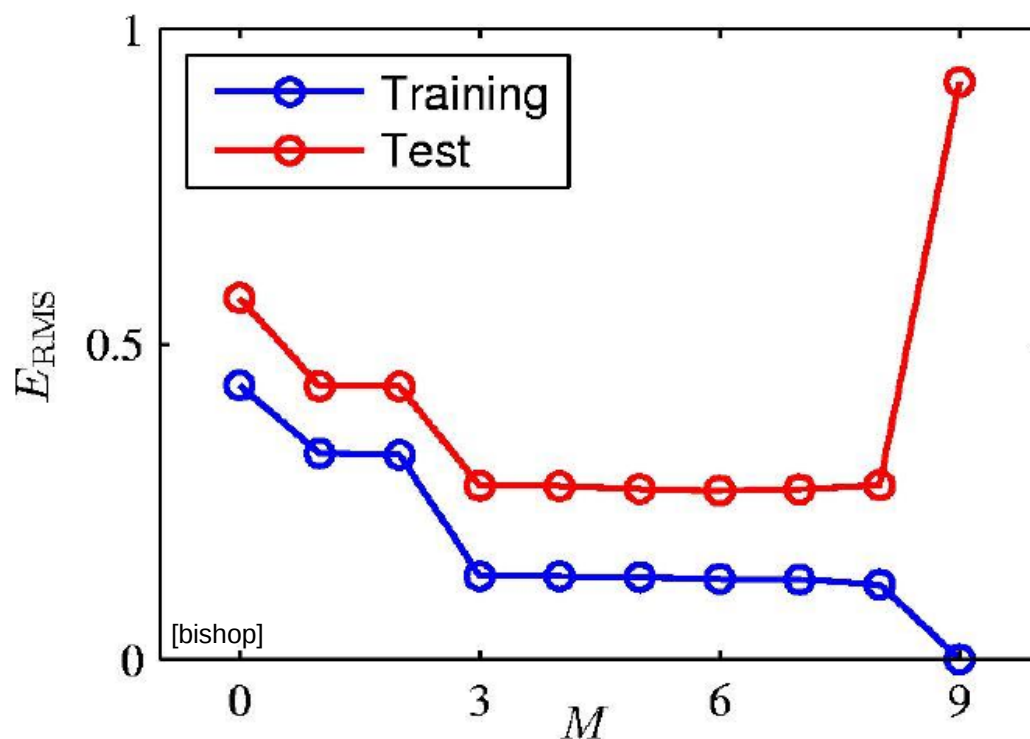


It is instructive to look at the **fitted weights** for various cases: when  $M$  increases the coefficient become **fine tuned** to data by developing large positive and negative values.

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43

## Root mean squared error (RMS)

$$E_{RMS} = \sqrt{\frac{1}{N} \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2} = \sqrt{\frac{E(\mathbf{w})}{N}}$$

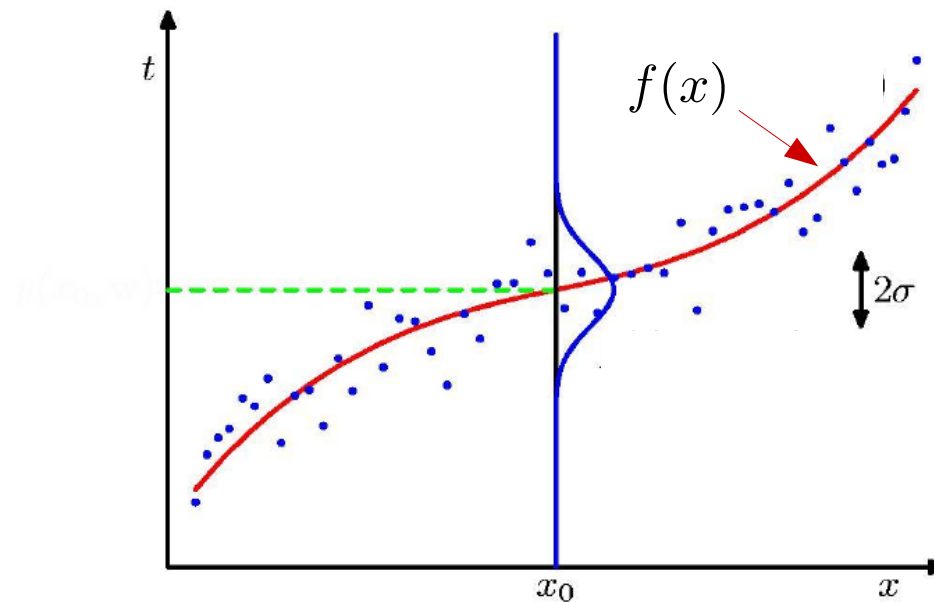


# The Bias-Variance decomposition

## Training dataset

- N observations of **feature**  $x = \{x_1, \dots, x_N\}$
- N **Target** values  $t = \{t_1, \dots, t_N\}$

We assume that  $t$  are distributed following a function:  $t_i = f(x_i) + \epsilon$



Noise  
(Mean 0,  
variance  $\sigma^2$ )

→ We want to find  $y(x)$  that approximate true function  $f(x)$

# The Bias-Variance decomposition (\*)

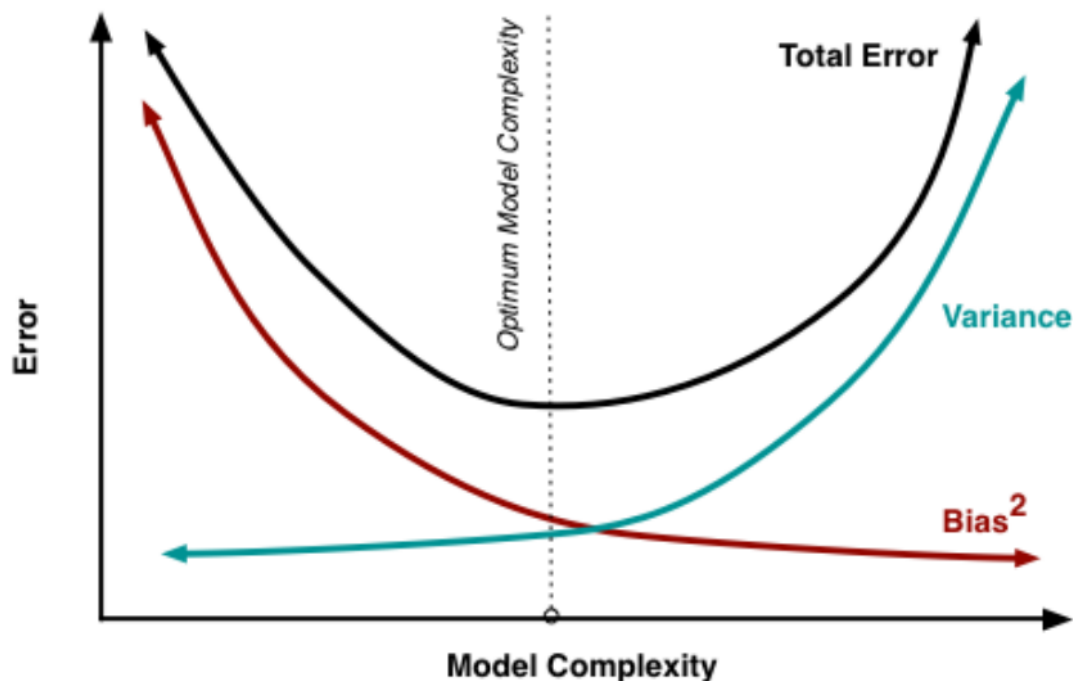
As before we determine  $y(x)$  by **minimizing**:  $\sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2$   
over the **training** dataset

The **expected error** for a **new test sample**  $\mathbf{x}$  can be decomposed as:

- Data **noise**: minimal error of the model
- **Bias** in the model: error caused by model assumptions
- **Variance** of model: how much  $y(x)$  depends on structure of data

$$\text{squared error on } y(x) = \boxed{\sigma^2} + \boxed{(\bar{y}(x) - f(x))^2} + \boxed{E[(y(x) - \bar{y}(x))^2]}$$

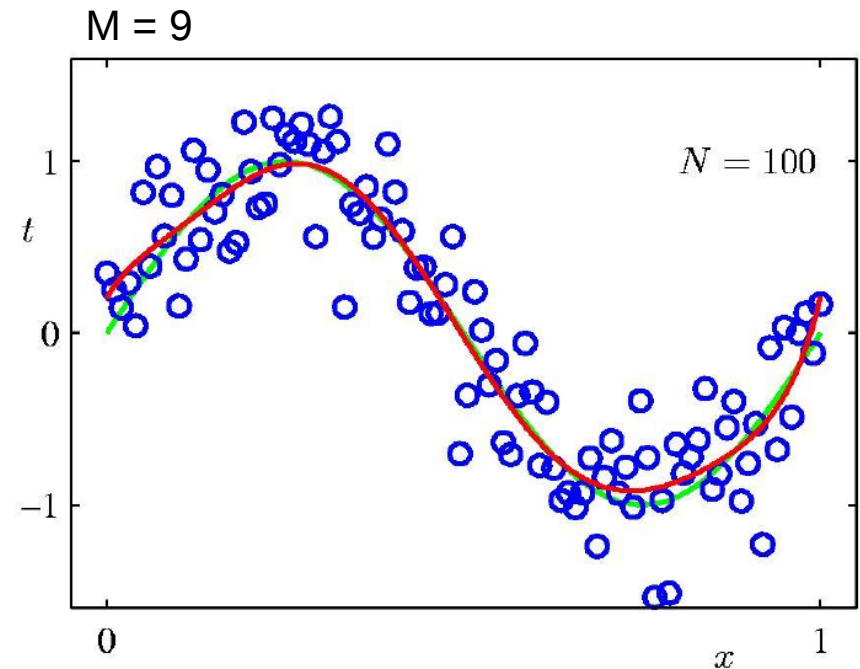
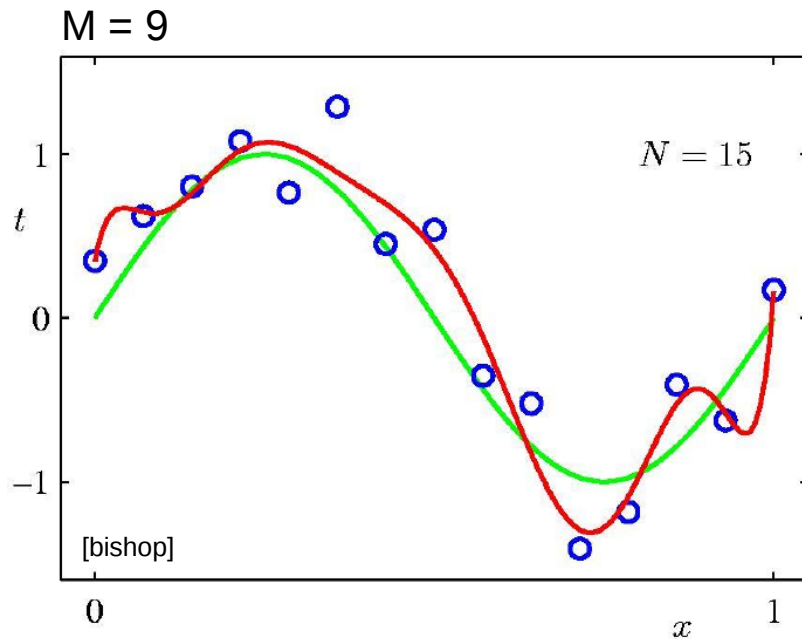
# The Bias-Variance decomposition



**Simple** models **under-fit**: deviate from data (high bias) but not influenced by structure of data (low variance)

**Complex** models **over-fit**: small deviation from data (low bias) but very sensitive to data fluctuations (high variance)

Overfitting really depends on **N** data and **M** parameters.



How can we constrain the fitted parameter into reasonable values ?

→ **Regularization** techniques can be a solution.

Add **penalization term** to error function in order to **constrain** parameters  $\mathbf{w}$ .

→ Simple penalization: **ridge regression** (L2 norm)

Constrains weight to be not too large .

$$\tilde{E}(\mathbf{w}) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2 + \lambda \|\mathbf{w}\|^2$$

where  $\|\mathbf{w}\|^2 = \mathbf{w}^T \mathbf{w} = w_0^2 + \dots + w_M^2$

and  $\lambda$  : parameter that governs the importance of regularization

## Other choices

- **Lasso** regression (L1 norm):  $\|\mathbf{w}\| = |w_0| + \dots + |w_M|$   
Reduce number of weights (set some of them to 0)
- **Elastic net**: L1 + L2 norm

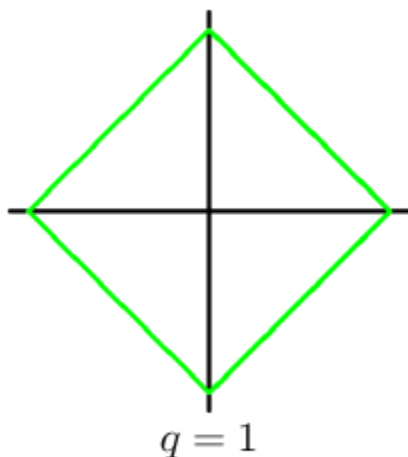
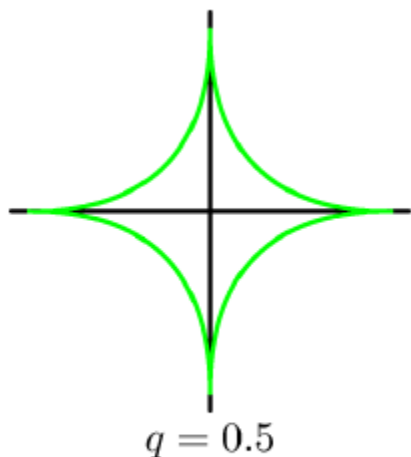


General regularization term is of the form:

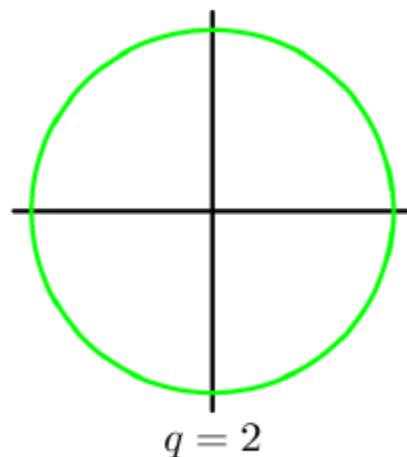
$$\tilde{E}(\mathbf{w}) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2 + \lambda \sum_{j=1}^M |w_j|^q$$

Minimizing this error function is equivalent to minimizing the unregularized sum-of-square error with the constraint

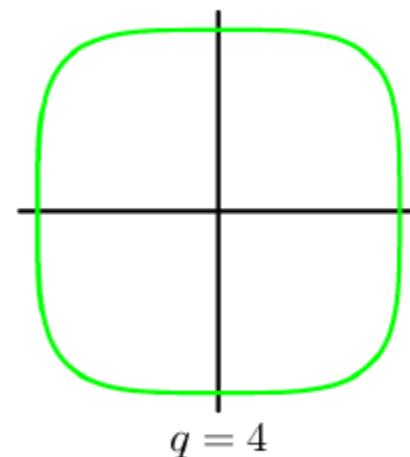
$$\sum_{j=1}^M |w_j|^q \leq \eta$$



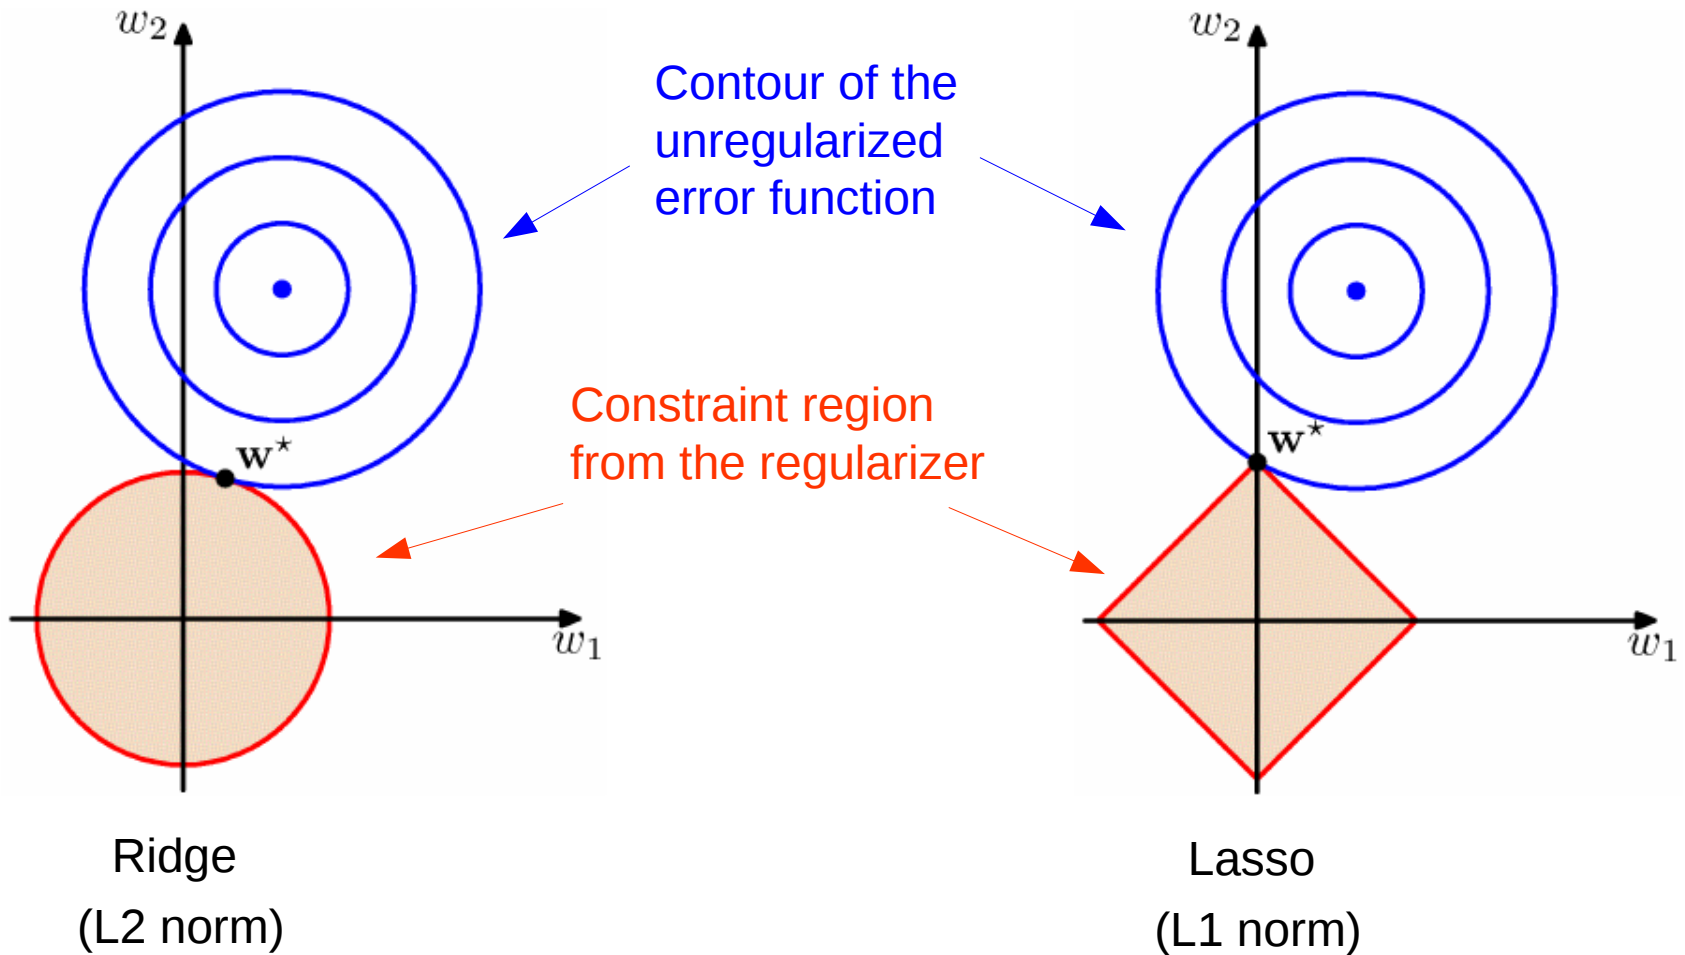
Lasso  
(L1 norm)



Ridge  
(L2 norm)

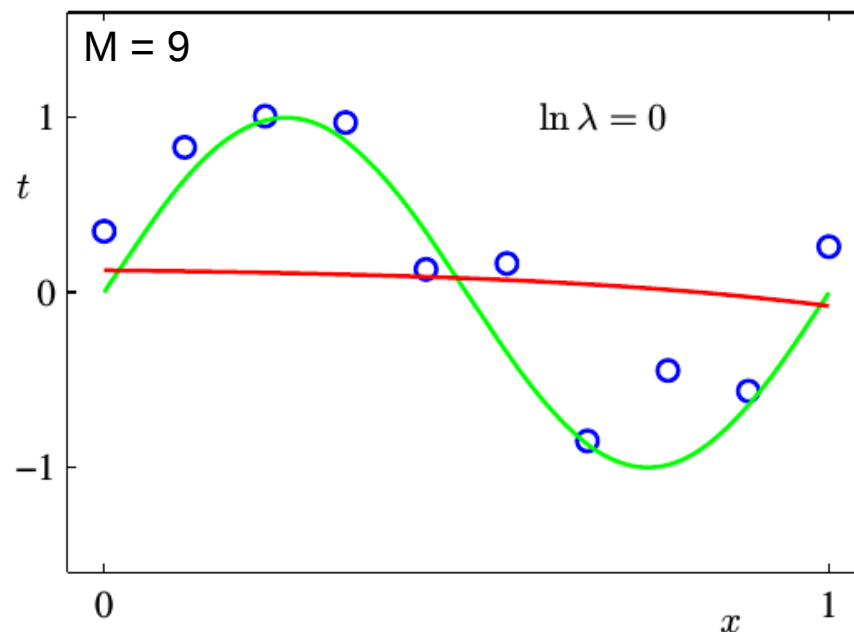
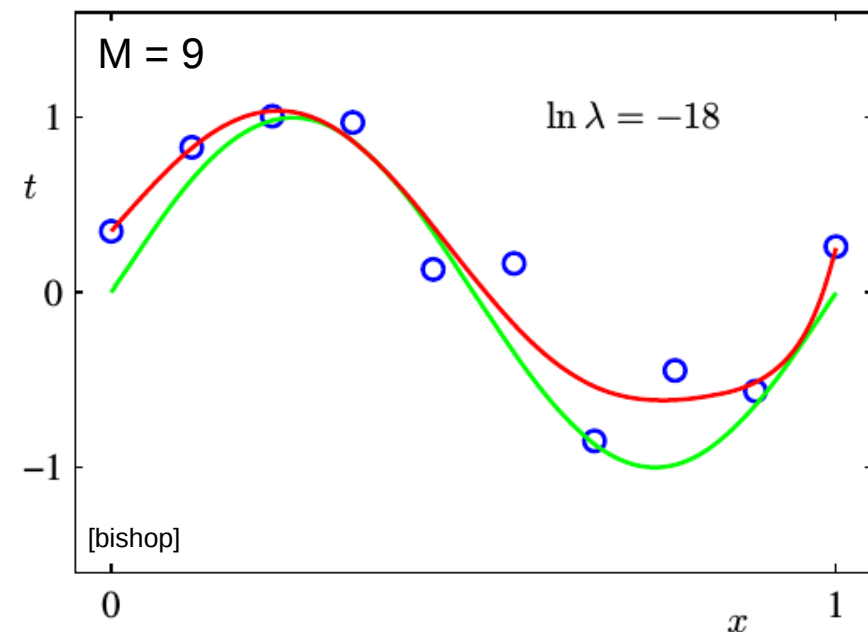


# Regularization



The optimum value for the parameter vector  $w$  is denoted by  $w^*$ .  
The lasso gives a sparse solution in which  $w_1^* = 0$ .

# Regularization



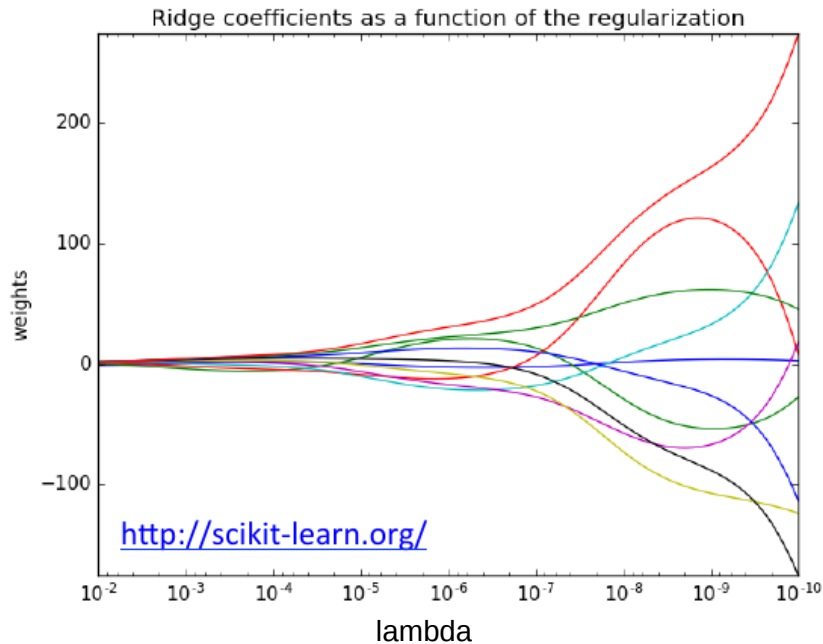
	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01

## Effect of L2 norm regularization

- $\ln \lambda = -\infty$  : no regularization
- $\ln \lambda = -18$  : suppressed overfitting
- $\ln \lambda = 0$  : fit too constrained

# Regularization

L2 norm:  $\lambda ||\mathbf{w}||^2$

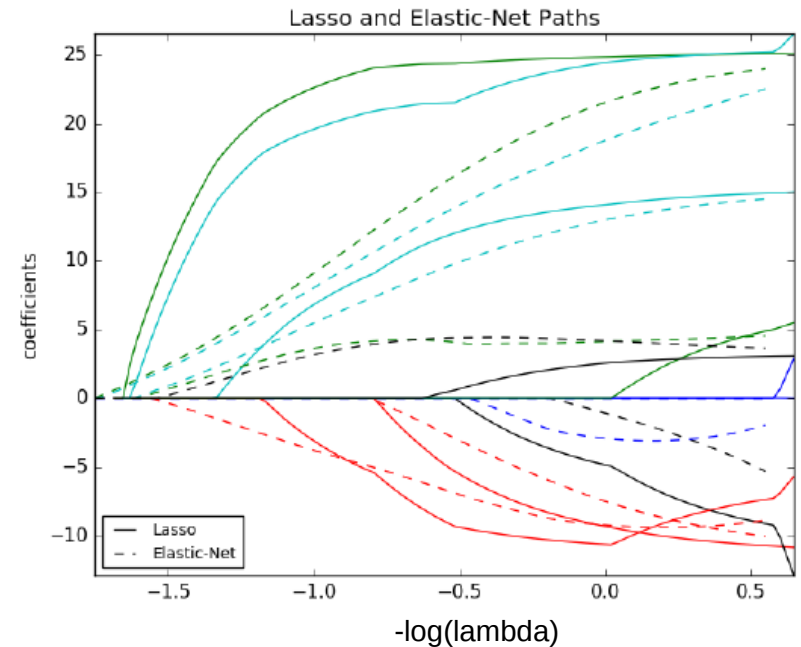


←

More constraints                      Less constraints

Affects value of coefficients  
(shrinkage)

L1 norm:  $\lambda ||\mathbf{w}||$

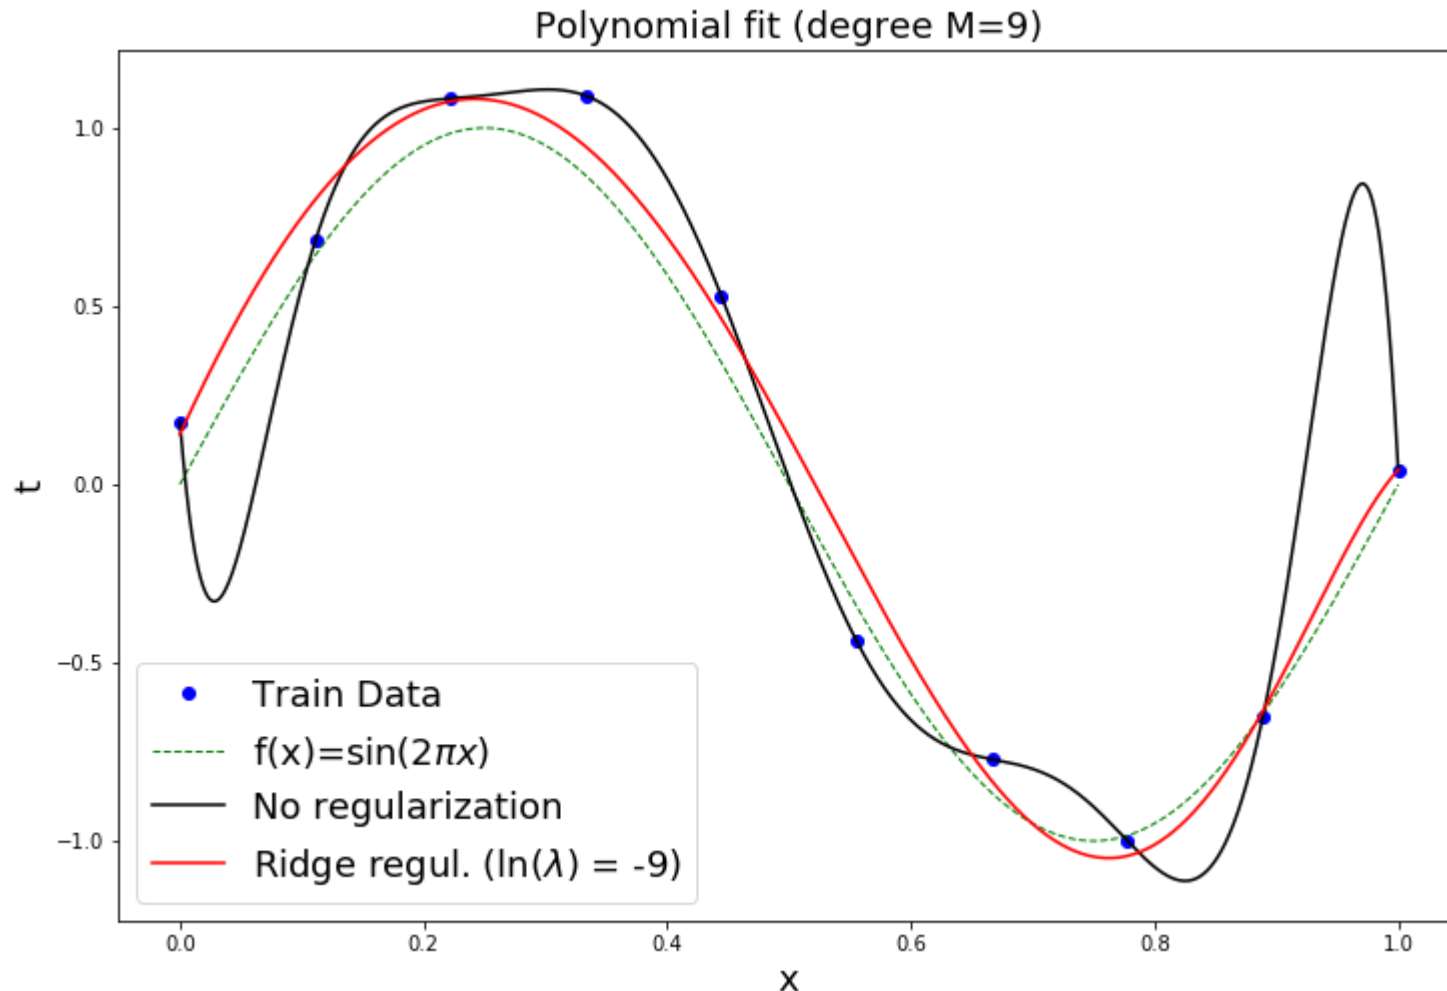


←

More constraints                      Less constraints

Affects number of coefficients  
(sparsity)

# Tutorial: Python code



<https://mybinder.org/v2/gh/judonini/MLcourses/master>



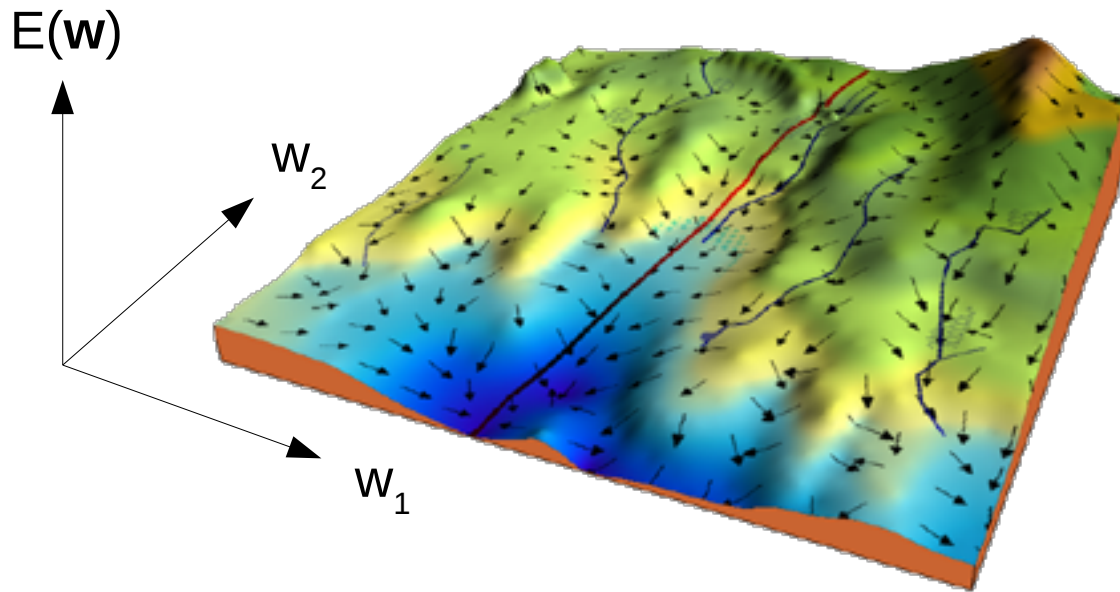
→ Polynomial-regression-regularization.ipynb



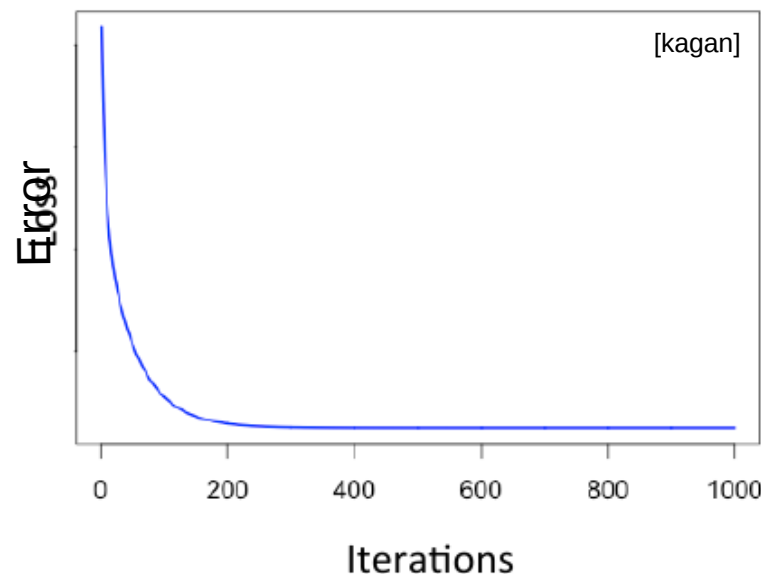
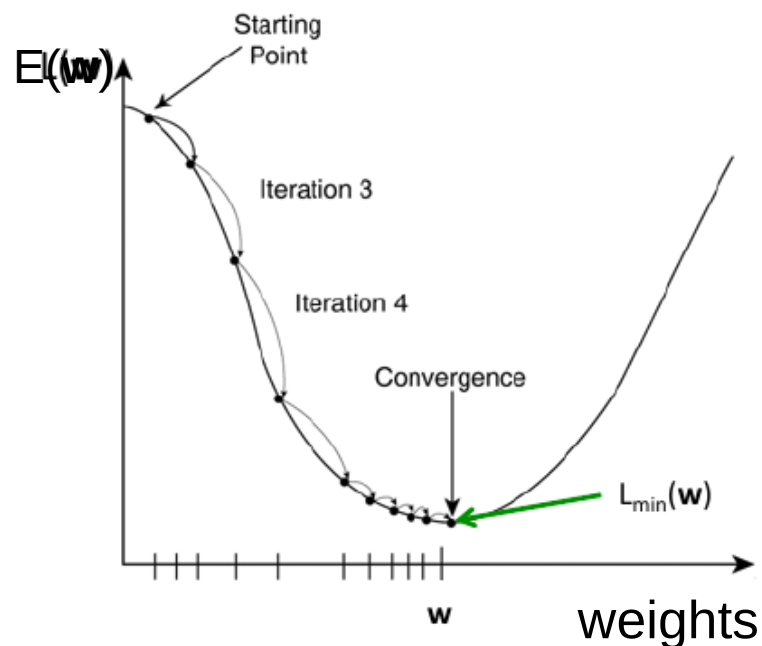
How can we **minimize** the error function for complex cases (ex: when there is no analytic solution) ?

→ **Solution: Gradient descent**

Iteratively move in the direction of steepest descent as defined by the negative of the gradient of the error function



Descend along the error function to find a (local) minimum:



Direction of descent:

→ (negative of the **gradient** of the error function) × (**learning rate**)



**Simple example:** fit N data points with linear function:  $y(x, \mathbf{w}) = w_0 + w_1 x$

**Error function and its derivatives**

$$E(w_0, w_1) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2 = \sum_{i=1}^N \{(w_0 + w_1 x_i) - t_i\}^2$$

$$\longrightarrow \begin{cases} \frac{\partial E(w_0, w_1)}{\partial w_0} = \sum_{i=1}^N 2 \{(w_0 + w_1 x_i) - t_i\} \\ \frac{\partial E(w_0, w_1)}{\partial w_1} = \sum_{i=1}^N 2 x_i \{(w_0 + w_1 x_i) - t_i\} \end{cases}$$

Iterative update **rule**:

$$\begin{aligned} w_0^{(k)} &\rightarrow w_0^{(k+1)} = w_0^{(k)} - \frac{\partial E(w_0, w_1)}{\partial w_0} \times \eta \\ w_1^{(k)} &\rightarrow w_1^{(k+1)} = w_1^{(k)} - \frac{\partial E(w_0, w_1)}{\partial w_1} \times \eta \end{aligned}$$

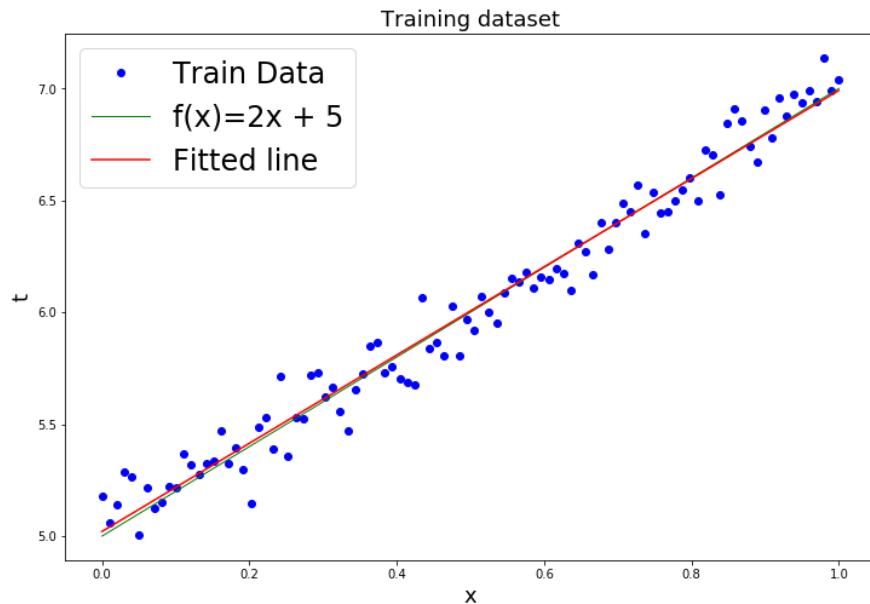
k: iteration number

$\eta$ : learning rate

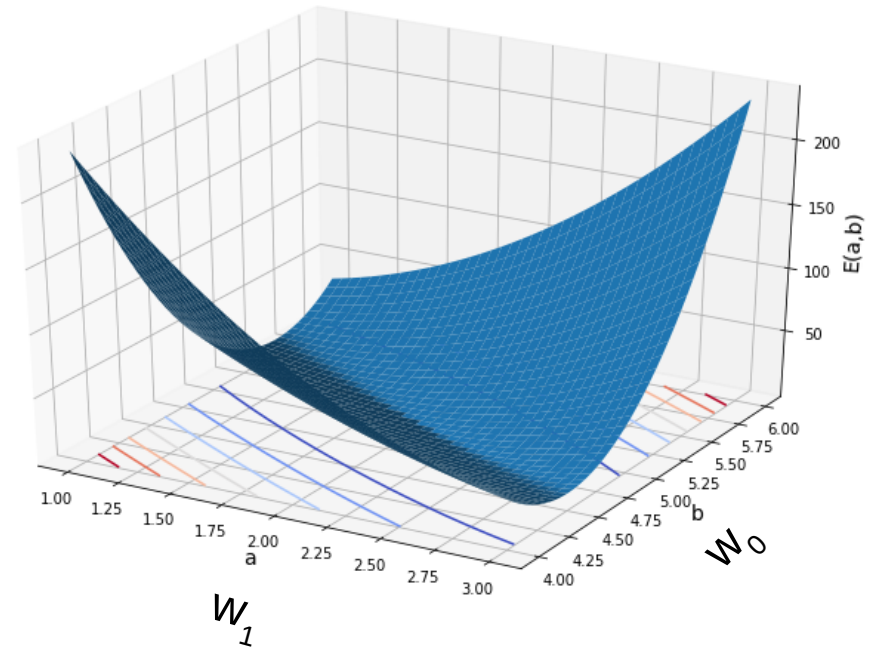
Repeat until convergence

# Gradient descent

Input data:  $\{x_i, t_i\}$



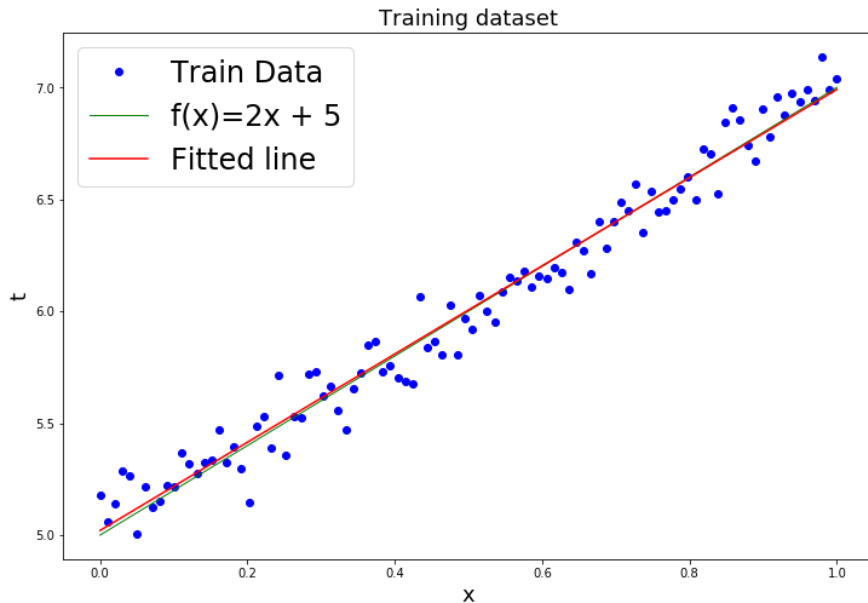
Error function



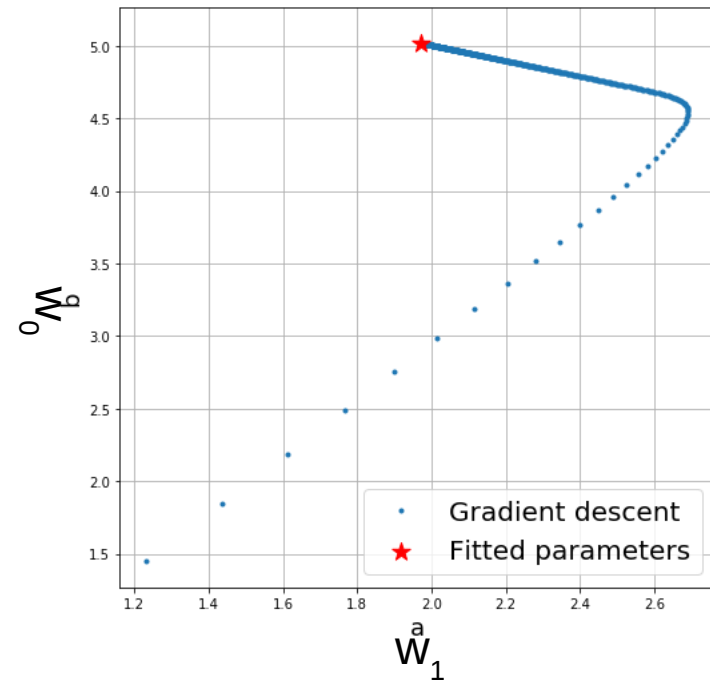
$$E(w_0, w_1) = \sum_{i=1}^N \{(w_0 + w_1 x_i) - t_i\}^2$$

# Gradient descent

Input data:  $\{x_i, t_i\}$



Gradient descent



$$w_0^{(k)} \rightarrow w_0^{(k+1)} = w_0^{(k)} - \frac{\partial E(w_0, w_1)}{\partial w_0} \times \eta$$

$$w_1^{(k)} \rightarrow w_1^{(k+1)} = w_1^{(k)} - \frac{\partial E(w_0, w_1)}{\partial w_1} \times \eta$$

1000 iterations


learning rate  $\eta = 0.05$

# Stochastic gradient descent

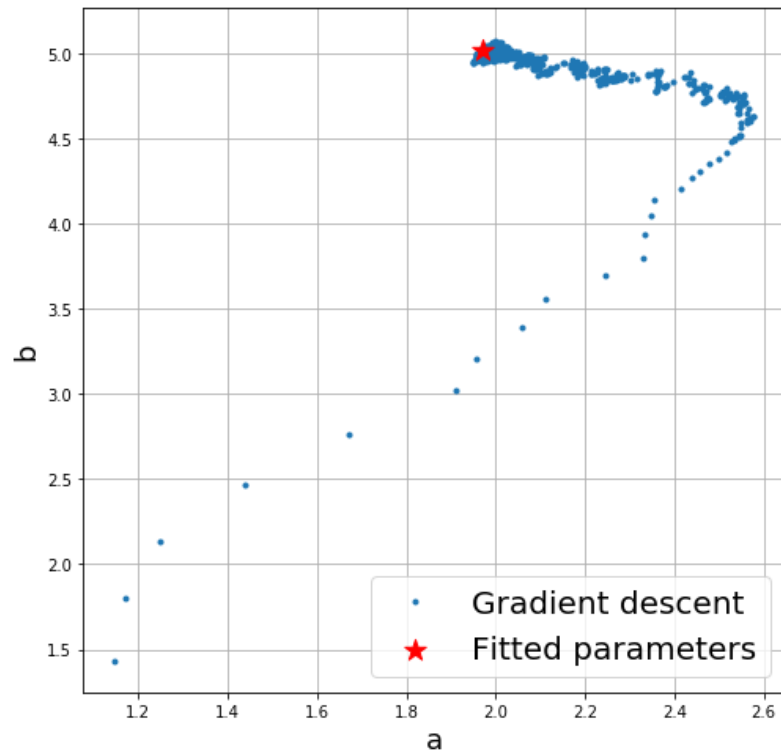
Gradient descent can be **computationally costly** for large N since the gradient is calculated over full training set.

→ **Solution: Stochastic gradient descent**

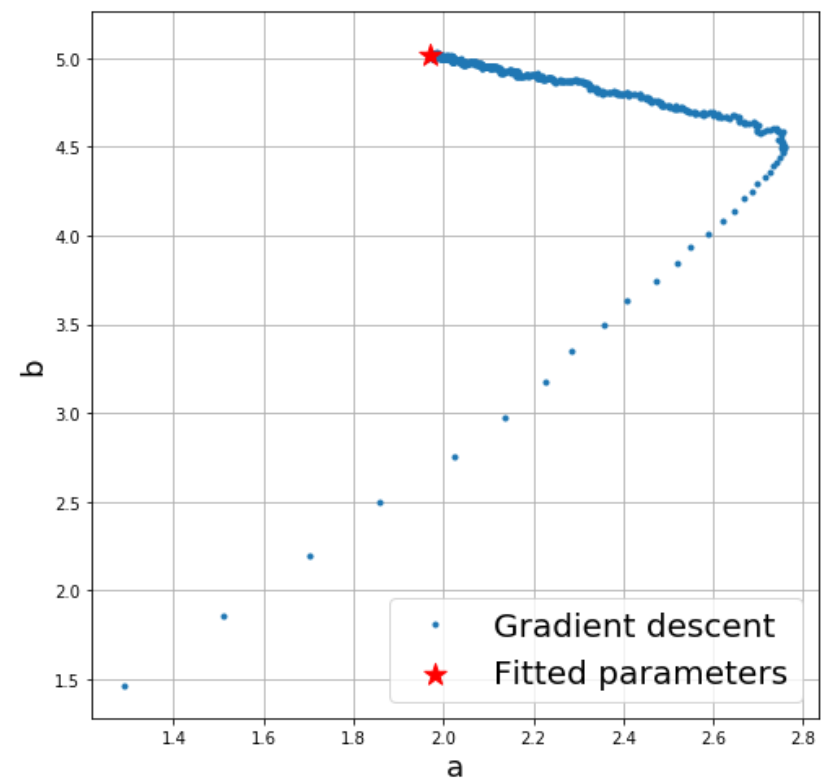
Compute gradient on a small **batch** of events (can be 1 event):

$$\begin{cases} \frac{\partial E(w_0, w_1)}{\partial w_0} = \sum_{i \in N} 2 \{ (w_0 + w_1 x_i) - t_i \} \\ \frac{\partial E(w_0, w_1)}{\partial w_1} = \sum_{i \in N} 2 x_i \{ (w_0 + w_1 x_i) - t_i \} \end{cases}$$


# Stochastic gradient descent



**Gradient** calculated on **1**  
(random) event at each step



**Gradient** calculated on **10**  
(random) events at each step



# Likelihood and regression (\*)

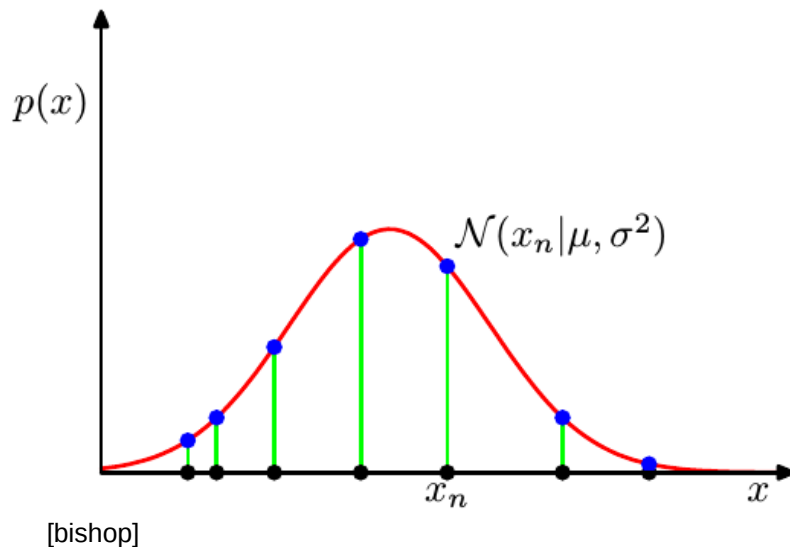
## Likelihood

Consider  $\mathbf{N}$  measurements of  $x$  distributed along a given probability law  $p(x)$ .

$$\mathbf{x} = (x_1, \dots, x_N)^\top$$

where values  $x_i$  are **independent and identically distributed** (i.i.d).

Ex: Normal (a.k.a Gaussian) law with 2 parameters: mean  $\mu$  and variance  $\sigma^2$



$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}$$

What is the *likelihood* of this set of measurements ?

Can we estimate  $\mu$  and  $\sigma$  given  $\mathbf{x}$  ?

# Likelihood and regression (\*)

## Likelihood and parameter estimation

Since the variables  $\mathbf{x}$  are i.i.d we can write the joint probability distribution, therefore the **likelihood** of the dataset, given  $\mu$  and  $\sigma$  is:

$$\mathcal{L}(\mu, \sigma^2; \mathbf{x}) \equiv p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \mathcal{N}(x_n|\mu, \sigma^2)$$

To **estimate**  $\mu$  and  $\sigma$  given  $\mathbf{x}$  one **maximizes**  $p$  w.r.t these parameters.  
In practice often maximize  $\ln(p)$  or minimize  $-\ln(p)$ .

$$\ln p(\mathbf{x}|\mu, \sigma^2) = -\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi)$$

$$\begin{cases} \frac{\partial(\ln p(\mathbf{x}|\mu, \sigma^2))}{\partial \mu} = 0 \rightarrow \mu_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N x_n \\ \frac{\partial(\ln p(\mathbf{x}|\mu, \sigma^2))}{\partial \sigma} = 0 \rightarrow \sigma_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N (x_n - \mu_{\text{ML}})^2 \end{cases}$$

### Expected values

$$\begin{aligned} \mathbb{E}[\mu_{\text{ML}}] &= \mu \\ \mathbb{E}[\sigma_{\text{ML}}^2] &= \left(\frac{N-1}{N}\right) \sigma^2 \end{aligned}$$



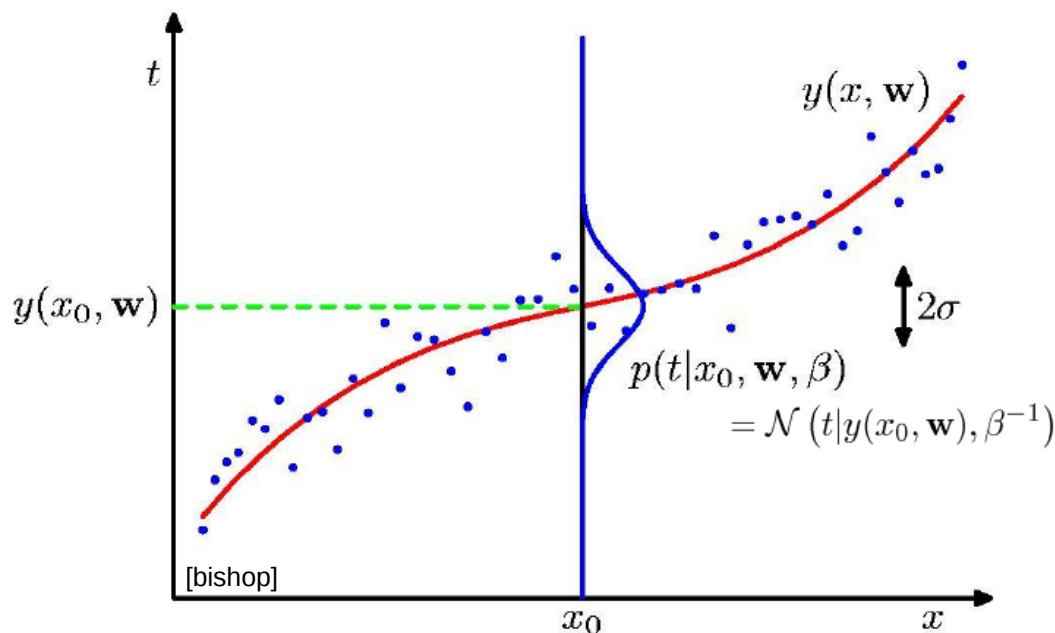
# Likelihood and regression (\*)

## Curve fitting with noise

Assume target variable in training dataset is subject to Gaussian noise

$$p(t|x, \mathbf{w}, \beta) = \mathcal{N}(t|y(x, \mathbf{w}), \beta^{-1})$$

where  $\beta = \frac{1}{\sigma^2}$  is a precision parameter.



# Likelihood and regression (\*)

## Predictive probabilistic model

By maximizing the likelihood on the training dataset we obtain a probabilistic predictive model for  $t$  (instead of a single point estimate):

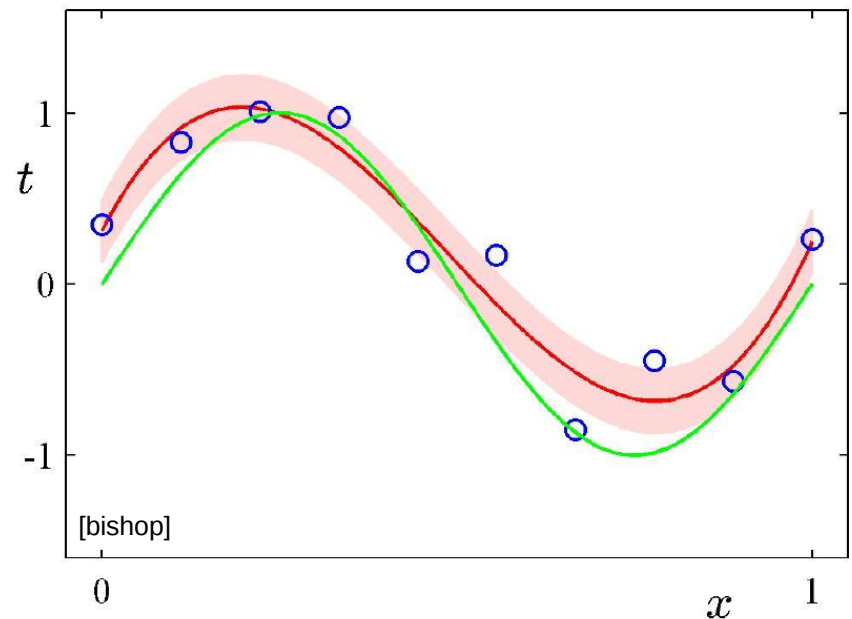
$$p(t|x, \mathbf{w}_{\text{ML}}, \beta_{\text{ML}}) = \mathcal{N}(t|y(x, \mathbf{w}_{\text{ML}}), \beta_{\text{ML}}^{-1})$$

where  $\mathbf{w}_{\text{ML}}$  is obtained by minimizing the sum of square error  $E(\mathbf{w})$

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

and  $\beta_{\text{ML}}$  is given by

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{n=1}^N \{y(x_n, \mathbf{w}_{\text{ML}}) - t_n\}^2$$



Exercise: show this !

# Likelihood and regression (\*)

Hints →

$$p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | y(x_n, \mathbf{w}), \beta^{-1})$$

$$\ln p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = - \underbrace{\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2}_{\beta E(\mathbf{w})} + \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi)$$

$\mathbf{w}_{\text{ML}}$  is determined by minimizing sum of square error  $E(\mathbf{w})$ , then:

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{n=1}^N \{y(x_n, \mathbf{w}_{\text{ML}}) - t_n\}^2$$

# Likelihood and regression (\*)

## Bayes theorem and likelihood



$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Posterior knowledge on theory

Likelihood of observing these data given a theory

Prior knowledge on theory

$$P(\text{theory}|\text{data}) = \frac{P(\text{data}|\text{theory})P(\text{theory})}{P(\text{data})}$$

Usually just a normalisation factor

# Likelihood and regression (\*)

## Bayesian approach

We assume that unknown parameters  $\mathbf{w}$  follow a Gaussian **prior** of the form:

$$p(\mathbf{w}|\alpha) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I}) = \left(\frac{\alpha}{2\pi}\right)^{(M+1)/2} \exp\left\{-\frac{\alpha}{2}\mathbf{w}^T\mathbf{w}\right\}$$

According to **Bayes Theorem** we have the **posterior**:

$$p(\mathbf{w}|\mathbf{x}, \mathbf{t}, \alpha, \beta) \propto p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta)p(\mathbf{w}|\alpha)$$

Taking  $-\ln(p) \rightarrow$  the maximum of the posterior is given by the minimum of:

$$\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \boxed{\frac{\alpha}{2} \mathbf{w}^T \mathbf{w}}$$

Regularization parameter with  $\lambda = \alpha/\beta$



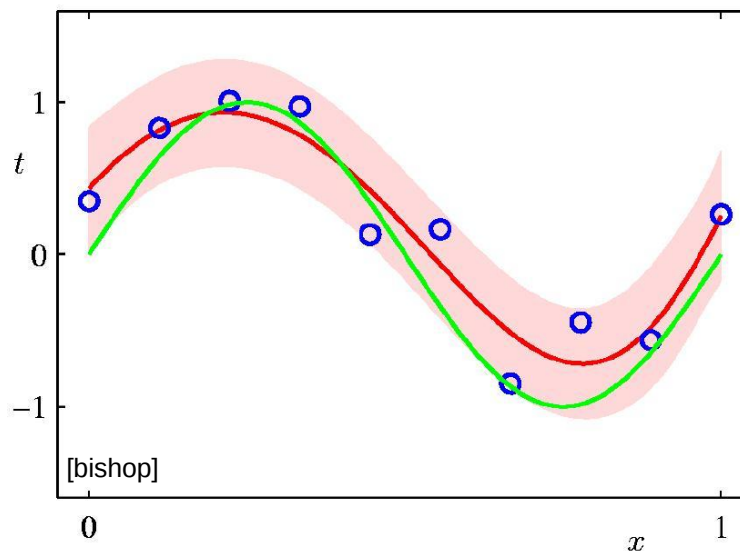
## Bayesian curve fitting

$$p(t|x, \mathbf{x}, \mathbf{t}) = \int p(t|x, \mathbf{w})p(\mathbf{w}|\mathbf{x}, \mathbf{t}) d\mathbf{w} = \mathcal{N}(t|m(x), s^2(x))$$

$$m(x) = \beta \phi(x)^T \mathbf{S} \sum_{n=1}^N \phi(x_n) t_n \quad s^2(x) = \beta^{-1} + \phi(x)^T \mathbf{S} \phi(x)$$

$$\mathbf{S}^{-1} = \alpha \mathbf{I} + \beta \sum_{n=1}^N \phi(x_n) \phi(x_n)^T \quad \phi(x_n) = (x_n^0, \dots, x_n^M)^T$$

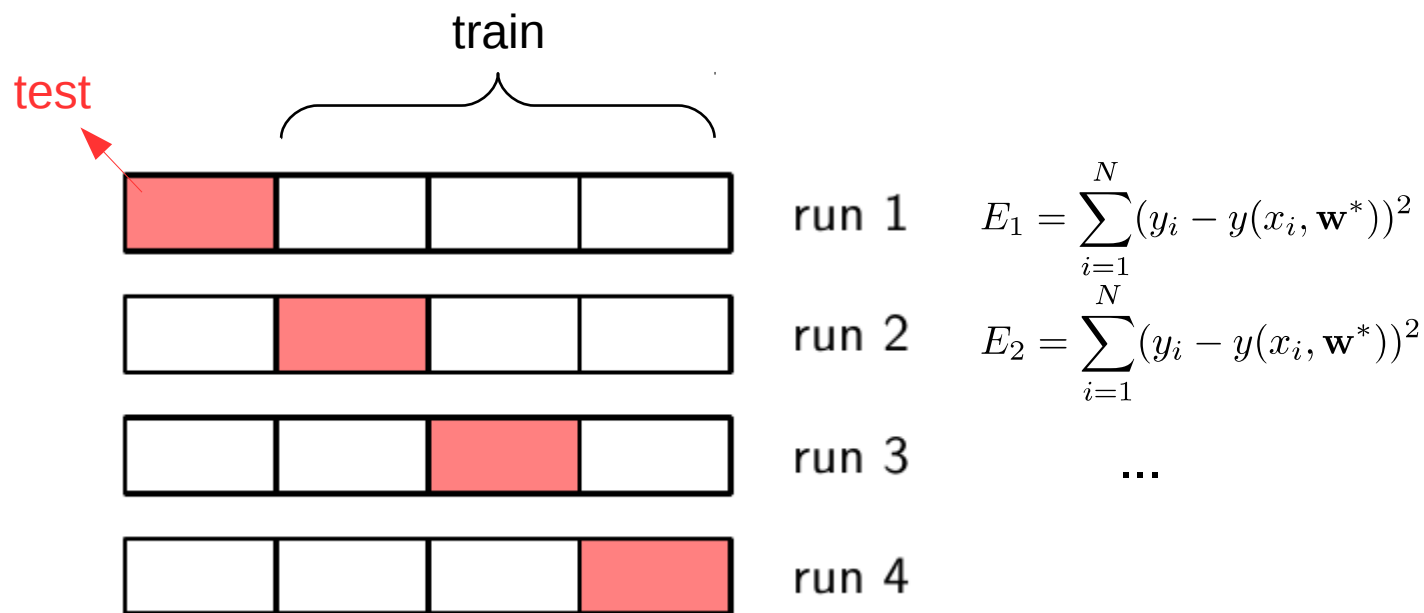
$$p(t|x, \mathbf{x}, \mathbf{t}) = \mathcal{N}(t|m(x), s^2(x))$$



## S-fold cross-validation

Divide data in S groups, use S-1 for training and test on left-over group

Rinse and repeat S times



**Cross-validation error:**  $CV = \frac{1}{S} \sum E_i$

Choose the set of parameters  $\mathbf{w}$  that give the smallest CV.

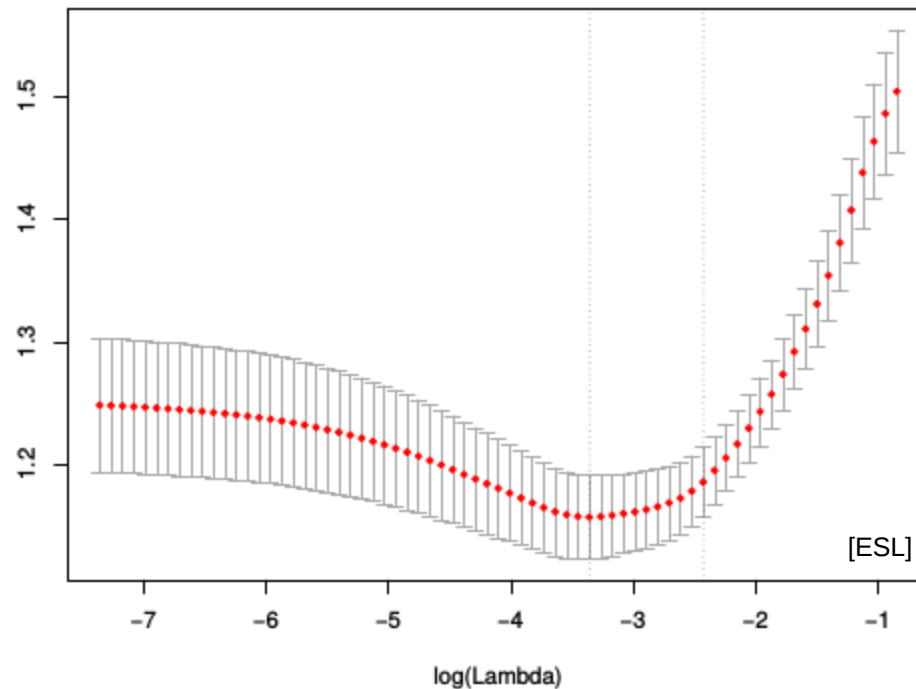
Drawback: can be very **time consuming** ...



## S-fold cross-validation

Divide data in  $S$  groups, use  $S-1$  for training and test on left-over group

Rinse and repeat  $S$  times



Cross-validation curve